

LV: Analysing Quality of Care data - resources

Catherine Blatter

2023-04-11

Table of contents

Preface	3
Necessary packages	3
Import dataset	3
Check the import	4
1 Duplicates	5
1.1 Different ‘types’ of duplicates	5
1.1.1 Fully duplicated entries	5
1.1.2 Fully distinct entries	6
1.1.3 Identify duplicates by counting the key (or key-index-pair)	6
2 Missings	8
2.1 Explore missings in sample data	8
2.2 Dataframe from slides	8
3 Visualising data	11
3.1 Prepare data for visualisation	11
3.2 Option 1: wide-to-long a.k.a stack	13
3.3 Option 2: Aggregation of values before plotting	15

Preface

This web-resource contains code that was talked about during class.

Necessary packages

If you do not have these installed (“Could not find package xyz”) you need to install the package.

```
# load necessary packages
# data preparation
library(readr)
library(tibble)
library(dplyr)

# visualization
library(ggplot2)

# missings
library(naniar)

# tables
library(kableExtra)
```

Import dataset

- By storing it into an object with the `zz_`-prefix, you can easier modify it afterwards without re-importing

```
# import the dataset using the {readr}-package
zz_dataset <- read_csv("data/2023-02-20_QoC_sample_dataset.csv",
                      show_col_types = FALSE)
```

Check the import

By looking at the top five rows (the “head” of the dataset):

```
head(zz_dataset, n = 5)
```

```
# A tibble: 5 x 6
  patient_ids sex    unit    adm          los disch
      <dbl> <chr> <chr> <date>      <dbl> <date>
1     13679 f     ward_B 2022-01-27    10 2022-02-06
2     15231 f     ward_A 2022-01-16     3 2022-01-19
3     10758 f     ward_A 2022-01-10    18 2022-01-28
4     10319 m     ward_B 2022-01-20     7 2022-01-27
5     17826 m     ward_A 2022-01-26     5 2022-01-31
```

- Seems the input worked as expected :relieved:
- the columns ‘sex’ and ‘unit’ should be a factor, not a character

```
# change 'sex' into factor
dataset <-
  zz_dataset %>%
  mutate(sex = as.factor(sex),
         unit = as.factor(unit))
```

... then again check the top five rows:

```
head(dataset, n = 5)
```

```
# A tibble: 5 x 6
  patient_ids sex    unit    adm          los disch
      <dbl> <fct> <fct> <date>      <dbl> <date>
1     13679 f     ward_B 2022-01-27    10 2022-02-06
2     15231 f     ward_A 2022-01-16     3 2022-01-19
3     10758 f     ward_A 2022-01-10    18 2022-01-28
4     10319 m     ward_B 2022-01-20     7 2022-01-27
5     17826 m     ward_A 2022-01-26     5 2022-01-31
```

1 Duplicates

1.1 Different ‘types’ of duplicates

- fully duplicated entries are ‘easy’ to catch but rare
- partially duplicated or mismatched entries are frequent but require more sophisticated checking

1.1.1 Fully duplicated entries

- `base::duplicated()` tests for fully duplicated entries (same row twice)
- several options to use:

```
# check numbers of duplicated rows
table(duplicated(dataset))
```

```
FALSE  TRUE
    99     1
```

- add a column `is_dupe` with information:

```
# adding column
dataset$is_dupe <- duplicated(dataset)
```

```
# check entry
head(dataset)
```

```
# A tibble: 6 x 7
  patient_ids sex    unit    adm          los disch      is_dupe
    <dbl> <fct> <chr>   <date>      <dbl> <date>    <lgl>
1    13679 f     ward_B 2022-01-27    10 2022-02-06 FALSE
2    15231 f     ward_A 2022-01-16     3 2022-01-19 FALSE
```

3	10758	f	ward_A	2022-01-10	18	2022-01-28	FALSE
4	10319	m	ward_B	2022-01-20	7	2022-01-27	FALSE
5	17826	m	ward_A	2022-01-26	5	2022-01-31	FALSE
6	17463	f	ward_A	2022-01-30	33	2022-03-04	FALSE

1.1.2 Fully distinct entries

- `dplyr::distinct()` tests for full distinct entries (same row twice)
- basically the opposite of `base::duplicated()`

1.1.3 Identify duplicates by counting the key (or key-index-pair)

```
# counting the number of rows per patient_ids
# using sort = T moves the highest number on top
count(dataset, patient_ids, sort = T)
```

```
# A tibble: 96 x 2
  patient_ids      n
    <dbl> <int>
1     10119      2
2     12335      2
3     16839      2
4     18098      2
5     10236      1
6     10319      1
7     10522      1
8     10647      1
9     10758      1
10    10875      1
# ... with 86 more rows
```

- there are several duplicated rows - although `duplicated()` only identified one fully duplicated entry
- this requires a further examination of these cases by identifying the “weird” ids and in a second create a subset (filtered) dataset, to check these cases:

```
# keep the weird ids
weird_ids <-
  count(dataset, patient_ids, sort = T) %>%
```

```

filter(n > 1 ) %>%
pull(patient_ids)

# filter the dataset by these values
dataset %>%
  filter(patient_ids %in% weird_ids) %>%
  arrange(patient_ids)


# A tibble: 8 x 7
  patient_ids sex    unit    adm        los disch    is_dupe
    <dbl> <fct> <chr>   <date>      <dbl> <date>    <lgl>
1     10119 f      ward_A 2022-01-08    27 2022-02-04 FALSE
2     10119 f      ward_A 2022-01-08    27 2022-02-04 TRUE
3     12335 m      ward_B 2022-01-15    21 2022-02-05 FALSE
4     12335 m      ward_B 2022-01-25    25 2022-02-19 FALSE
5     16839 m      ward_A 2022-01-08    16 2022-01-24 FALSE
6     16839 f      ward_B 2022-01-08    NA NA      FALSE
7     18098 f      ward_B 2022-01-13    16 2022-01-29 FALSE
8     18098 f      ward_B 2022-01-13    NA NA      FALSE

```

2 Missings

2.1 Explore missings in sample data

```
naniar::vis_miss(dataset)
```



- There are only few missings
- NA's in disch are linked to NA's in los

2.2 Dataframe from slides

- by using the command `set.seed(1234)` you should be able to reproduce the same dataframe as in the slides


```
# set seed for reproducibility
set.seed(1234)

# create a sample-df with 100 rows and
# 3 variables
df_miss <-
  tibble(id = 1:100,
          var1 = sample(
            c(1:4, NA_real_), 100, T),
          var2 = sample(
            c(1:4, NA_real_), 100, T))


# first 4 entries
head(df_miss, 4)
```

```
# A tibble: 4 x 3
      id  var1  var2
<int> <dbl> <dbl>
1     1     4     3
2     2     2     1
3     3    NA     3
4     4     4     4
```

```
# quick summary of proportional missings
naniar::miss_prop_summary(df_miss)
```

```
# A tibble: 1 x 3
      df  var  case
<dbl> <dbl> <dbl>
1 0.127 0.667 0.34
```

```
naniar::vis_miss(df_miss)
```



3 Visualising data

3.1 Prepare data for visualisation

Consider the example from class where four people (A, B, C, D) were asked about their tool-preferences for accessing their e-banking account:

- Formally, this dataframe is in ‘tidy’ format and the unit of analysis is the `person_id`

Now we represent the same information in a barchart, to compare the frequencies:


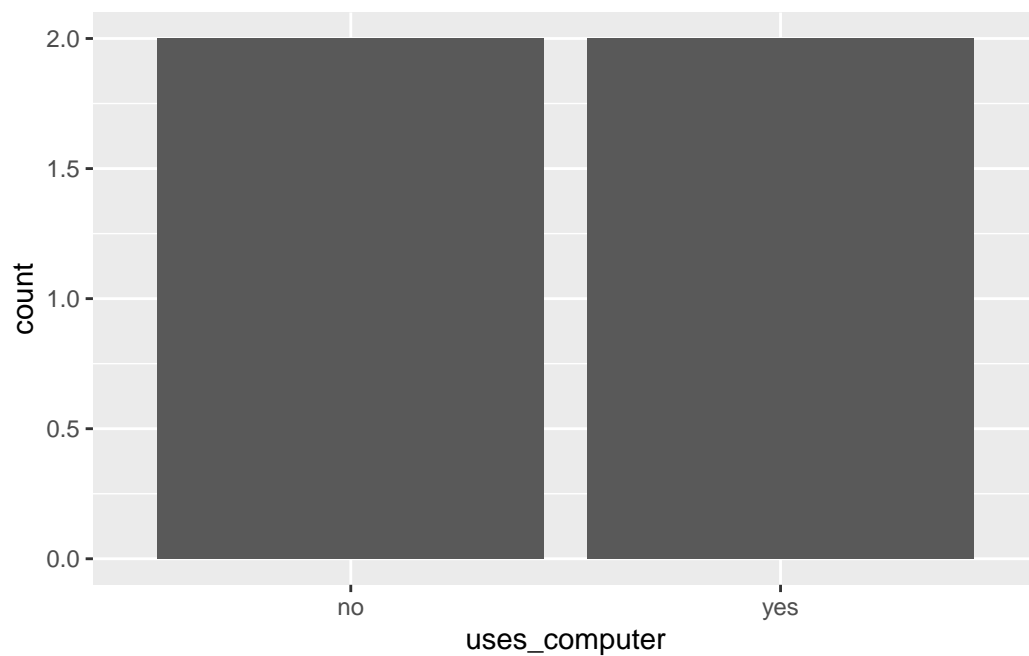


Table 3.1: Tidy dataframe - wide format

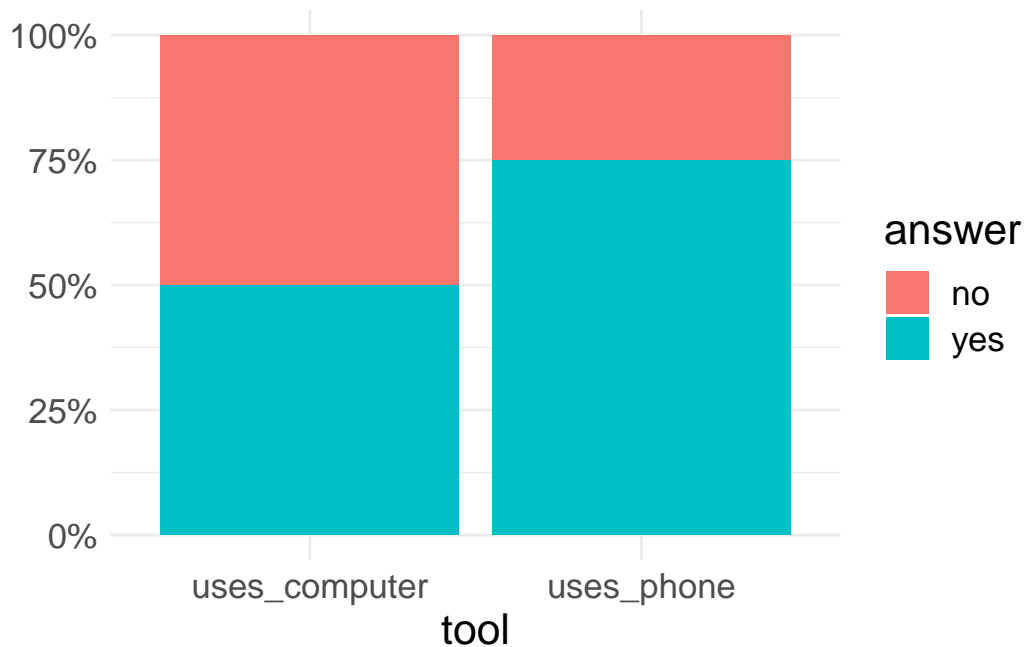
person_id	uses_phone	uses_computer
A	yes	no
B	yes	yes
C	yes	no
D	no	yes

But how did we actually input the data? If we try this we only get the variable `uses_computer` on the x-axis (with the values no and yes):

```
ggplot(data = wldf, aes(x = uses_computer)) +  
  geom_bar()
```



When we go back to the original plot:



- we see that we want a variable `tool` on the x-axis, that has the values 'uses_computer' and 'uses_phone'
- we want a variable `answer` that has the values 'no' and 'yes'

(- we possibly want a variable `percentage` on the y-axis, that has numeric values according to the proportions)

- In summary, the format of the data needs to change in order to plot

3.2 Option 1: wide-to-long a.k.a stack

In this option we use `tidyr::pivot_longer()` to create a dataframe in long format:


```
# we want to stack the columns uses_computer and uses_phone
long_df <-
  wldf %>%
  pivot_longer(cols = c(uses_computer, uses_phone),
               names_to = "tool",
               values_to = "answer")

# print the df
long_df
```

```
# A tibble: 8 x 3
  person_id tool      answer
  <chr>      <chr>    <chr>
1 A        uses_computer no
2 A        uses_phone  yes
3 B        uses_computer yes
4 B        uses_phone  yes
5 C        uses_computer no
6 C        uses_phone  yes
7 D        uses_computer yes
8 D        uses_phone  no
```

- this dataframe now has the variable `tool` and `answer` for the ggplot2-code using `geom_bar(position = "fill")` :


```
ggplot(long_df, aes(x = tool, fill = answer)) +
  geom_bar(position = "fill")
```



- of course, this needed some further 'layout-work' to improve the labeling of scales and the removal of the grey-background

```
ggplot(long_df, aes(x = tool, fill = answer)) +
  geom_bar(position = "fill") +
```

```
scale_y_continuous(name = NULL,
                   labels = scales::percent) +
theme_minimal() +
theme(text = element_text(size = 16))
```



3.3 Option 2: Aggregation of values before plotting

Another option is to aggregate (or summarise) values manually, before we pass them to a plotting function (the example above is not really a good example for this), but the code might be handy if you have more variables than 2 for this situation.

In the example, we'd need to have three variables as indicated above: `tool` (for the x-axis), `answer` for the fill and `percent` for the percentage. We can create this table with:

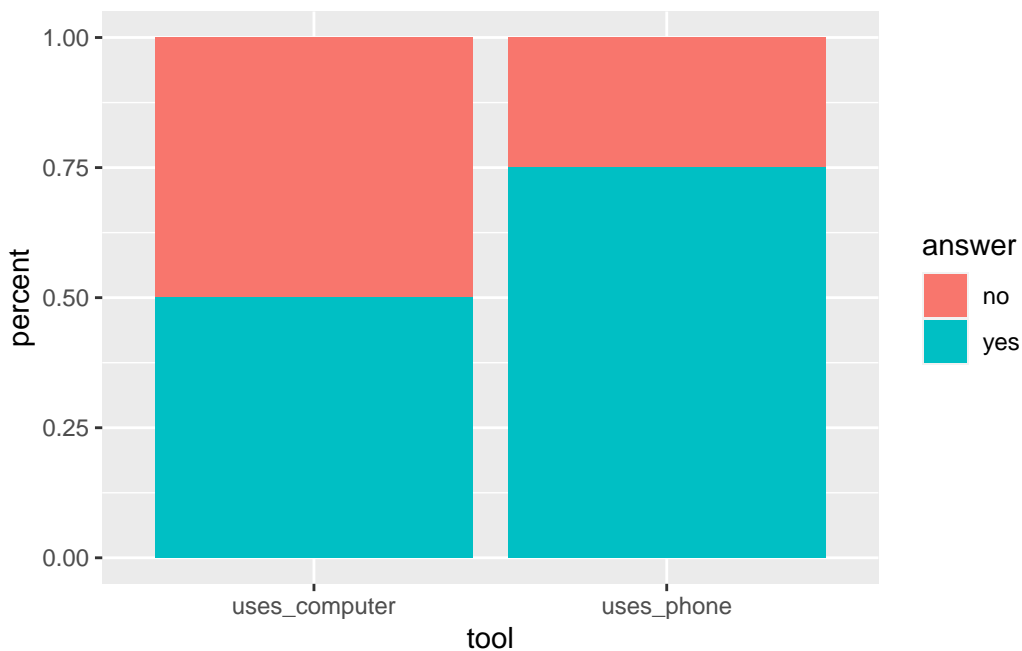
```
agg_data <-
  wldf |>
  select(uses_computer, uses_phone) |>
  as.list() |>
  map_dfr(~janitor::tabyl(.x), .id = "tool") |>
  rename(answer = .x)
```

```
# print agg_data
agg_data
```

	tool	answer	n	percent
	uses_computer	no	2	0.50
	uses_computer	yes	2	0.50
	uses_phone	no	1	0.25
	uses_phone	yes	3	0.75

And then we can plot it:

```
ggplot(data = agg_data, aes(x = tool, y = percent, fill = answer)) +  
  geom_col(position = "stack")
```



- Careful: instead of using `geom_bar(position="fill")` without a y-variable, here we pass the `percent`-value to the y-axis and use `geom_col(position="stack")` for the `answer`-variable

An finally, improve the layout:

```
ggplot(data = agg_data, aes(x = tool, y = percent, fill = answer)) +  
  geom_col(position = "stack") +
```



```
scale_y_continuous(name = NULL,  
                   labels = scales::percent) +  
theme_minimal() +  
theme(text = element_text(size = 16))
```

